

Experiment - 3

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud

- 1) Create 3 EC-2 instances with all running on Amazon Linux as OS with inbound SSH allowed.
- 2) To efficient run kubernetes cluster, select instance type of at least t2.medium as kubernetes recommends at least 2 vCPU to run smoothly

EC2 > Instances > Launch an instance

Launch an instance

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags

Name

worker

Add additional tags

Application and OS Images (Amazon Machine Image)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Q Search our full catalog including 1000s of application and OS images

Recents

Quick Start

Amazon Linux

macOS

Ubuntu

Windows

Red Hat

SUSE Linux

Browse more AMIs

Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Amazon Linux 2023 AMI

ami-0182f573e66f89c85 (64-bit x86), uefi-preferred / ami-0b947c5d5516fa06e (64-bit Arm), uefi

Virtualization: hvm ENA enabled: true Root device type: ebs

Free tier eligible

Summary

Number of instances

3

When launching more than 1 instance, consider EC2 Auto Scaling

Software Image (AMI)

Amazon Linux 2023 AMI 2023.5.2...read more

ami-0182f573e66f89c85

Virtual server type (instance type)

t2.medium

Firewall (security group)

New security group

Storage (volumes)

1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

Cancel

Launch instance

Review commands

Instance type

Instance type

t2.medium

Family: t2 - 2 vCPU 4 GiB Memory Current generation: true

On-Demand Linux base pricing: 0.0464 USD per Hour

On-Demand RHEL base pricing: 0.0752 USD per Hour

On-Demand Windows base pricing: 0.0664 USD per Hour

On-Demand SUSE base pricing: 0.1464 USD per Hour

All generations

Compare instance types

Key pair (login)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

brijesh

Create new key pair

Network settings

Network

vpc-0531204c9e29f6332

Subnet

No preference (Default subnet in any availability zone)

Auto-assign public IP

Enable

Additional charges apply when outside of free tier allowance

Firewall (security groups)

Create security group

Select existing security group

We'll create a new security group called 'launch-wizard-8' with the following rules:

Allow SSH traffic from

Anywhere

Summary

Number of instances

3

When launching more than 1 instance, consider EC2 Auto Scaling

Software Image (AMI)

Amazon Linux 2023 AMI 2023.5.2...read more

ami-0182f573e66f89c85

Virtual server type (instance type)

t2.medium

Firewall (security group)

New security group

Storage (volumes)

1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

Cancel

Launch instance

Review commands

3) Three instance are ready - master, worker1, and worker2.

Instances (1/9) Info

Last updated less than a minute ago

Refresh

Connect

Instance state ▾

Actions ▾

Launch instances ▾

All states ▾

< 1 >

⚙️

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP
<input type="checkbox"/>	master	i-076c3bb1d7c953b83	Running	t2.medium	Initializing	View alarms	us-east-1a	ec2-34-...
<input type="checkbox"/>	worker1	i-0e6c44ebfb4ca3afa	Running	t2.medium	Initializing	View alarms	us-east-1a	ec2-18-...
<input checked="" type="checkbox"/>	worker2	i-0d83b162ba0948d34	Running	t2.medium	Initializing	View alarms	us-east-1a	ec2-18-...

4) Connect the instances to the local terminal using the SSH client.

5) Run the following commands on all the machines.

Install Docker

a) `sudo yum install docker -y`

```

ec2-user@ip-172-31-25-225 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:06:08 ago on Thu Sep 12 13:45:25 2024.
Dependencies resolved.

Package                               Architecture      Version           Repository        Size
=====
Installing:
docker                               x86_64            25.0.6-1.amzn2023.0.2  amazonlinux      44 M
Installing dependencies:
containerd                           x86_64            1.7.20-1.amzn2023.0.1  amazonlinux      35 M
iptables-libs                        x86_64            1.8.8-3.amzn2023.0.2  amazonlinux      401 k
iptables-nft                         x86_64            1.8.8-3.amzn2023.0.2  amazonlinux      183 k
libcgroup                            x86_64            3.0-1.amzn2023.0.1    amazonlinux       75 k
libnetfilter_conntrack               x86_64            1.0.8-2.amzn2023.0.2  amazonlinux       58 k
libnftnl                             x86_64            1.0.1-19.amzn2023.0.2  amazonlinux       30 k
libnftnl                             x86_64            1.2.2-2.amzn2023.0.2  amazonlinux       84 k
pigpio                               x86_64            2.5-1.amzn2023.0.3    amazonlinux       83 k
runc                                  x86_64            1.1.13-1.amzn2023.0.1  amazonlinux      3.2 M

Transaction Summary
Install 10 Packages

Total download size: 84 M
Installed size: 317 M
Downloading Packages:
(1/10): iptables-libs-1.8.8-3.amzn2023.0.2.x86_64.rpm 3.6 MB/s | 401 kB 00:00

```

b) Then, configure cgroup in a daemon.json file by using following commands. This allows kubernetes to manage host more efficiently -

- `cd /etc/docker`

Run the scripts below -

```
cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
"storage-driver": "overlay2"
}
EOF
```

```
[ec2-user@ip-172-31-18-9 ~]$ cd /etc/docker
[ec2-user@ip-172-31-18-9 docker]$
```

```
[ec2-user@ip-172-31-18-9 docker]$ cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
"storage-driver": "overlay2"
}
EOF
{
"exec-opts": ["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
"storage-driver": "overlay2"
}
[ec2-user@ip-172-31-18-9 docker]$
```

c) After configuring restart docker service service :

- sudo systemctl enable docker
- sudo systemctl daemon-reload
- sudo systemctl restart docker
- docker -v

```
[ec2-user@ip-172-31-18-9 docker]$ sudo systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-18-9 docker]$
```

```
[ec2-user@ip-172-31-18-9 docker]$ sudo systemctl daemon-reload
sudo systemctl restart docker
docker -v
Docker version 25.0.5, build 5dc9bcc
[ec2-user@ip-172-31-18-9 docker]$
```

Install Kubernetes

- a) SELinux needs to be disabled before configuring kubelet to avoid interference with kubernetes api server

- sudo setenforce 0
- sudo sed -i 's/^SELINUX=enforcing\$/SELINUX=permissive/' /etc/selinux/config
- Add kubernetes repository (paste in terminal)

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
```

```
[ec2-user@ip-172-31-18-9 docker]$ sudo setenforce 0
[ec2-user@ip-172-31-18-9 docker]$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[ec2-user@ip-172-31-18-9 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[ec2-user@ip-172-31-18-9 docker]$
```

- b) sudo yum update

```
[ec2-user@ip-172-31-25-225 docker]$ sudo yum update
Kubernetes
Dependencies resolved.
Nothing to do.
Complete!
```

c) `sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes`

```
[ec2-user@ip-172-31-18-9 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Last metadata expiration check: 0:00:54 ago on Thu Sep 12 14:50:28 2024.
Dependencies resolved.
=====
Package                                Architecture      Version           Repository        Size
-----
Installing:
kubeadm                                x86_64            1.30.5-150500.1.1  kubernetes        10 M
kubectl                                x86_64            1.30.5-150500.1.1  kubernetes        10 M
kubelet                                x86_64            1.30.5-150500.1.1  kubernetes        17 M
Installing dependencies:
conntrack-tools                       x86_64            1.4.6-2.amzn2023.0.2  amazonlinux        208 k
cri-tools                             x86_64            1.30.1-150500.1.1  kubernetes         8.6 M
kubernetes-cni                        x86_64            1.4.0-150500.1.1  kubernetes         6.7 M
libnetfilter_cthelper                 x86_64            1.0.0-21.amzn2023.0.2  amazonlinux         24 k
libnetfilter_cttimeout                x86_64            1.0.0-19.amzn2023.0.2  amazonlinux         24 k
libnetfilter_queue                    x86_64            1.0.5-2.amzn2023.0.2  amazonlinux         30 k
Transaction Summary
-----
Install 9 Packages

Total download size: 53 M
Installed size: 292 M
Downloading Packages:
(1/9): conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64.rpm 3.1 MB/s | 208 kB 00:00
(2/9): libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64.rpm 332 kB/s | 24 kB 00:00
```

d) After installing Kubernetes, we need to configure internet options to allow bridging.

- `sudo swapoff -a`
- `echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf`
- `sudo sysctl -p`

```
[ec2-user@ip-172-31-25-225 docker]$ sudo swapoff -a
[ec2-user@ip-172-31-25-225 docker]$ echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
net.bridge.bridge-nf-call-iptables=1
[ec2-user@ip-172-31-25-225 docker]$ sudo sysctl -p
net.bridge.bridge-nf-call-iptables = 1
[ec2-user@ip-172-31-25-225 docker]$
```

6) To perform only on Master machine

a) Initialize kubernetes by typing below command

`sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all`

```
net.bridge.bridge-nf-call-iptables = 1
[ec2-user@ip-172-31-18-9 docker]$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all
I0912 14:55:34.563710 30027 version.go:256] remote version is much newer: v1.31.0; falling back to: stable-1.30
[init] Using Kubernetes version: v1.30.4
[preflight] Running pre-flight checks
[WARNING FileExisting-socat]: socat not found in system path
[WARNING FileExisting-tc]: tc not found in system path
[WARNING Service-Kubelet]: kubelet service is not enabled, please run 'systemctl enable kubelet.service'
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
W0912 14:55:34.706553 30027 checks.go:844] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.9" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-18-9.ec2.internal kubernetes kubernetes.default kubernetes.default.svc kubernetes.default.svc.cluster.local] and IPs [10.96.0.1 172.31.18.9]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
```

b) Copy this join link

`kubeadm join 172.31.22.128:6443 --token 2nzclk.1ek0i93tsqnednb9 \`

`--discovery-token-ca-cert-hash`

`sha256:e7c55b0579b7e928431704c459e9c9c521c4af034e3d346f3418e1afc672928d`

c) Run the below command -

- `mkdir -p $HOME/.kube`
- `sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`
- `sudo chown $(id -u):$(id -g) $HOME/.kube/config`

d) Then, add a common networking plugin called flannel file as mentioned in the code.
 kubectl apply -f
<https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

```
[ec2-user@ip-172-31-22-128 docker]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
[ec2-user@ip-172-31-22-128 docker]$ kubectl get pods
No resources found in default namespace.
```

- e) Check the created pod using this command
- kubectl get pods

7) To perform on both worker machine -

- a) sudo yum install iproute-tc socat -y

```
[ec2-user@ip-172-31-27-40 ~]$ sudo yum install iproute-tc socat -y
Last metadata expiration check: 0:11:39 ago on Sat Sep 14 12:44:32 2024.
Dependencies resolved.
===== Package Architecture
=====Installing:
iproute-tc      x86_64      5.10.0-2.amzn2023.0.5      amazonlinux      455 k
socat           x86_64      1.7.4.2-1.amzn2023.0.2     amazonlinux      303 k

Transaction Summary
=====Install 2 Packages

Total download size: 758 k
Installed size: 2.0 M
Downloading Packages:
(1/2): socat-1.7.4.2-1.amzn2023.0.2.x86_64.rpm 4.5 MB/s | 303 kB 00:00
(2/2): iproute-tc-5.10.0-2.amzn2023.0.5.x86_64.rpm 6.1 MB/s | 455 kB 00:00
-----Total

Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
  Preparing      : 1/1
  Installing     : socat-1.7.4.2-1.amzn2023.0.2.x86_64 1/2
  Installing     : iproute-tc-5.10.0-2.amzn2023.0.5.x86_64 2/2
  Running scriptlet: iproute-tc-5.10.0-2.amzn2023.0.5.x86_64 2/2
  Verifying      : iproute-tc-5.10.0-2.amzn2023.0.5.x86_64 1/2
  Verifying      : socat-1.7.4.2-1.amzn2023.0.2.x86_64 2/2

Installed:
  iproute-tc-5.10.0-2.amzn2023.0.5.x86_64 socat-1.7.4.2-1.amzn2023.0.2.x86_64

Complete!
```

- b) sudo systemctl enable kubelet

```
[ec2-user@ip-172-31-17-38 ~]$ sudo systemctl enable kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service.
```

- c) `sudo systemctl restart kubelet`
- d) `kubeadm join 172.31.22.128:6443 --token 2nzcl.1ek0i93tsqnednb9 \`
`--discovery-token-ca-cert-hash sha256:e7c55b0579b7e928431704c459e9c9c521c4af034e3d346f3418e1afc672928d`

```
[ec2-user@ip-172-31-17-38 ~]$ sudo kubeadm join 172.31.22.128:6443 --token 2nzcl.1ek0i93tsqnednb9 --discovery-token-ca-cert-hash sha256:e7c55b0579b7e928431704c459e9c9c521c4af034e3d346f3418e1afc672928d
[preflight] Running pre-flight checks
error execution phase preflight: couldn't validate the identity of the API Server: failed to request the cluster-info ConfigMap: Get "https://172.31.22.128:6443/api/v1/namespaces/kube-public/configmaps/cluster-info?timeout=10s": context deadline exceeded
To see the stack trace of this error execute with --v=5 or higher
```

If it gives error or refusing connection, restart the master server using `sudo systemctl restart kubelet` and try to connect.

With the help of command the worker nodes are connected master node and is ready to do task assigned by master node.

Conclusion -

In this experiment, we connected master nodes from Kubernetes to the worker nodes successfully. First, we created the instances and connected with Kubernetes, while installing and configuring, there were some packages, that were needed to be installed separately. Even while connecting the nodes, error occurs and hence, system needs to be restarted and connected properly..