

Experiment - 4

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy your first Kubernetes Application.

1) Create an EC2 instance and allow SSH traffic to connect.

Instances (1/3)

Info

Last updated

1 minute ago

Refresh

Connect

Instance state ▾

Actions ▾

Launch instances ▾

Find Instance by attribute or tag (case-sensitive)

All states ▾

Instance state = running X






Clear filters

<

1

>

⚙

<input checked="" type="checkbox"/>	Name  ▾	Instance ID	Instance state ▾	Instance type ▾	Status check	Alarm status	Availability Zone ▾	Public IP
<input checked="" type="checkbox"/>	Node	i-0ef27b61aa74379bd	 Running  	t2.medium	 2/2 checks pass View alarms +		us-east-1a	ec2-54-175-68-222

2) Select the instance and click on Connect and go to Connect.

EC2 > Instances > i-0ef27b61aa74379bd > Connect to instance

Connect to instance Info


Connect to your instance i-0ef27b61aa74379bd (Node) using any of these options


EC2 Instance Connect

Session Manager

SSH client

EC2 serial console


**Port 22 (SSH) is open to all IPv4 addresses**
Port 22 (SSH) is currently open to all IPv4 addresses, indicated by 0.0.0.0/0 in the inbound rule in [your security group](#). For increased security, consider restricting access to only the EC2 Instance Connect service IP addresses for your Region: 18.206.107.24/29. [Learn more](#).

Instance ID
 i-0ef27b61aa74379bd (Node)


Connection Type

☒ **Connect using EC2 Instance Connect**
Connect using the EC2 Instance Connect browser-based client, with a public IPv4 address.

☐ **Connect using EC2 Instance Connect Endpoint**
Connect using the EC2 Instance Connect browser-based client, with a private IPv4 address and a VPC endpoint.

Public IPv4 address
 54.175.68.222

Username
Enter the username defined in the AMI used to launch the instance. If you didn't define a custom username, use the default username, ec2-user.

 **Note:** In most cases, the default username, ec2-user, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel

Connect

```
{
  "exec-opts": ["native.cgroup"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
```

```
[ec2-user@ip-172-31-30-107 ~]$ cd /etc/docker
[ec2-user@ip-172-31-30-107 docker]$ cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
```

6) Run the following command after this:

```
sudo systemctl enable docker
```

```
sudo systemctl daemon-reload
```

```
sudo systemctl restart docker
```

```
[ec2-user@ip-172-31-30-107 docker]$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
```

7) Check for installation of docker by docker -v command.

```
[ec2-user@ip-172-31-30-107 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc
```

Install Kubernetes

8) Disable SELinux before configuring kubelet

- sudo setenforce 0
- sudo sed -i 's/^SELINUX=enforcing\$/SELINUX=permissive/' /etc/selinux/config

```
[ec2-user@ip-172-31-30-107 docker]$ sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

9) Add kubernetes repository

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
```

```
[kubernetes]
```

```
name=Kubernetes
```

```
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
```

```
enabled=1
```

```

gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.k
ey
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF

```

10) Run the commands to update and install kubernetes packages

```
sudo yum update
```

```
sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
```

```
[ec2-user@ip-172-31-31-241 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Last metadata expiration check: 0:00:46 ago on Sat Sep 14 16:39:07 2024.
Dependencies resolved.
=====
Package                               Architecture      Version            Repository          Size
-----
Installing:
kubeadm                               x86_64            1.31.1-150500.1.1  kubernetes          11 M
kubectl                               x86_64            1.31.1-150500.1.1  kubernetes          11 M
kubelet                               x86_64            1.31.1-150500.1.1  kubernetes          15 M
Installing dependencies:
container-tools                       x86_64            1.4.6-2.amzn2023.0.2  amazonlinux        208 k
cri-tools                             x86_64            1.31.1-150500.1.1  kubernetes          6.9 M
kubernetes-cni                       x86_64            1.5.1-150500.1.1  kubernetes          7.1 M
libnetfilter_cthelper                x86_64            1.0.0-21.amzn2023.0.2  amazonlinux         24 k
libnetfilter_cttimeout                x86_64            1.0.0-19.amzn2023.0.2  amazonlinux         24 k
libnetfilter_queue                   x86_64            1.0.5-2.amzn2023.0.2  amazonlinux         30 k
Transaction Summary
-----
Install 9 Packages
Total download size: 51 M
Installed size: 269 M
Downloading Packages:
(1/9) kubelet-1.31.1-150500.1.1.x86_64.rpm: 15 MB/s | 0 B/s | 00:00
(2/9) kubeadm-1.31.1-150500.1.1.x86_64.rpm: 11 MB/s | 0 B/s | 00:00
(3/9) kubectl-1.31.1-150500.1.1.x86_64.rpm: 11 MB/s | 0 B/s | 00:00
(4/9) container-tools-1.4.6-2.amzn2023.0.2.x86_64.rpm: 11 MB/s | 0 B/s | 00:00
(5/9) cri-tools-1.31.1-150500.1.1.x86_64.rpm: 11 MB/s | 0 B/s | 00:00
(6/9) libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64.rpm: 11 MB/s | 0 B/s | 00:00
(7/9) libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64.rpm: 11 MB/s | 0 B/s | 00:00
(8/9) libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64.rpm: 11 MB/s | 0 B/s | 00:00
(9/9) kubernetes-cni-1.5.1-150500.1.1.x86_64.rpm: 11 MB/s | 0 B/s | 00:00
Total 15 MB/s | 0 B/s | 00:00

```

11) Configure internet options to allow bridging

```
sudo swapoff -a
```

```
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
```

```
sudo sysctl -p
```

```
[ec2-user@ip-172-31-31-241 docker]$ sudo swapoff -a
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
sudo sysctl -p
net.bridge.bridge-nf-call-iptables=1
net.bridge.bridge-nf-call-iptables = 1

```

12) Initialize the kubecuster

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
```

```
[ec2-user@ip-172-31-31-241 docker]$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[WARNING FileExisting-socat]: socat not found in system path
[WARNING FileExisting-tc]: tc not found in system path
[WARNING Service-Kubelet]: kubelet service is not enabled, please run 'systemctl enable kubelet.service'
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W0914 16:43:42.619918 27909 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.10" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-31-241.ec2.internal kubernetes kubernetes.default kubernetes.default.svc kubernetes.default.svc.cluster.local] and IPs [10.96.0.1 172.31.31.241]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-31-241.ec2.internal localhost] and IPs [172.31.31.241 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-31-241.ec2.internal localhost] and IPs [172.31.31.241 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key

```

13) Copy the Join Command-

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.31.241:6443 --token x6zuwa.bafjhtm4fqxp4yi8 \
--discovery-token-ca-cert-hash sha256:ad0a2b0eb8318975f42be705f750f923adc01bd102ebd7d93401df81429ef5a5
```

```
kubeadm join 172.31.31.241:6443 --token x6zuwa.bafjhtm4fqxp4yi8 \
--discovery-token-ca-cert-hash
```

```
sha256:ad0a2b0eb8318975f42be705f750f923adc01bd102ebd7d93401df81429ef5a5
```

14) Run the following commands

```
[ec2-user@ip-172-31-31-241 docker]$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
cp: overwrite '/home/ec2-user/.kube/config'? yes
```

15) Add a common network plugin called Flannel as mentioned in the code below:

```
kubectl apply -f
```

```
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

```
ec2-user@ip-172-31-20-245 dockekubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
ec2-user@ip-172-31-20-245 docker]$
```

16) Check for creation of pods

```
[ec2-user@ip-172-31-20-245 ~]$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	Pending	0	80s

17) To change the state from pending to running, use the following command kubectl describe pod nginx

This command will help to describe the pods it gives reason for failure as it shows the untolerated taints which need to be untainted.

```
Containers:
  nginx:
    Image:          nginx:1.14.2
    Port:           80/TCP
    Host Port:      0/TCP
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-dmncs (ro)
Conditions:
  Type              Status
  PodScheduled      False
Volumes:
  kube-api-access-dmncs:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:    kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:      true
QoS Class:          BestEffort
Node-Selectors:      <none>
Tolerations:        node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                    node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type      Reason              Age   From          Message
  ----      -
Warning    FailedScheduling    2m47s default-scheduler  0/1 nodes are available: 1 node(s) had intolerated taint (node-role.kubernetes.io/control-plane: ). Preemption: 0/1 nodes are available: 1 Preemption is not helpful for scheduling.
```

```
[ec2-user@ip-172-31-20-245 docker]$ kubectl taint nodes ip-172-31-20-245.ec2.internal node-role.kubernetes.io/control-plane-
node/ip-172-31-20-245.ec2.internal untainted
```

18) Check the status of pods.

```
[ec2-user@ip-172-31-20-245 docker]$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	4m3s

```
[ec2-user@ip-172-31-20-245 docker]$ kubectl port-forward nginx 8081:80
Forwarding from 127.0.0.1:8081 -> 80
Forwarding from [::1]:8081 -> 80
```

Conclusion:

We successfully configured Kubernetes environment on Amazon Linux EC2 instance. We installed Docker and adjusted its settings to use systemd for cgroup management. Then, we installed Kubernetes. We disabled SELinux, and added the Kubernetes repository. Then we installed the necessary components. Then added a network plugin and deployed on Nginx server. We also addressed issues related to pod scheduling and port forwarding, ensuring the Nginx pod.